

Task No: 6

Date:

Implement CURD operations on Casandra

Tools: Casandra, Docker desktop

CO3

## AIM:

To install and configure the apache Cassandra in windows operating system, and to perform the CURD operations.

## PROCEDURE:

1. Install Docker desktop, Java 8 or latest version, and Python
2. Visit the Apache Cassandra website at <https://cassandra.apache.org/download/> to download the latest stable release of Cassandra.
3. Start the Docker engine and type the following commands in the command prompt
4. Check the Docker Version

**docker --version**

5. Creating a Cluster of two Cassandra nodes. Name the nodes Cassandra-1 and Cassandra-2  
**docker run --name cassandra-1 -d cassandra:latest**
6. Run the second Container, Cassandra-2, and link it to Cassandra-1.  
**docker run --name cassandra-2 -d --link cassandra-1:cassandra cassandra:latest**
7. Check the status of our node using the nodetool utility.  
**docker exec -it cassandra-1 nodetool status**  
**docker exec -it cassandra-2 nodetool status**
8. Use a Command Line Interface (CQLSH) to interact with the Cluster to create databases and tables using bash command  
**docker exec -it cassandra-1 bash -c 'cqlsh'**
9. Create the keyspace and tables, then execute the CURD operations – create, update, read, delete.
10. Execute various no-sql queries and find the results.
11. Stop the cassandra server.  
**docker stop cassandra-1**  
**docker stop cassandra-2**

## Syntax:

### Create Keyspace:

A keyspace is like RDBMS database which contains column families, indexes, user defined types, data center awareness, strategy used in keyspace, replication factor, etc.

```
CREATE KEYSPACE <identifier> WITH <properties>
```

Or

```
Create keyspace KeyspaceName with replication={'class':strategy name,  
'replication_factor': No of replications on different nodes}
```

### Alter Keyspace:

```
ALTER KEYSPACE <identifier> WITH <properties>
```

Or

```
ALTER KEYSPACE "KeySpace Name"
```

```
WITH replication = {'class': 'Strategy name', 'replication_factor' : 'No.Of replicas'};
```

### Drop Keyspace:

```
DROP keyspace KeyspaceName ;
```

### Create Table Operation:

Syntax of Create Operation-

```
CREATE (TABLE | COLUMNFAMILY) <tablename>  
(<column-definition> , <column-definition>)  
(WITH <option> AND <option>)
```

### CRUD Operatons:

```
CREATE KEYSPACE test_cassandra WITH replication = {'class':'SimpleStrategy' , 'replication_factor' :  
1};
```

```
Use test_cassandra ;
```

```
CREATE TABLE student(  
student_id int PRIMARY KEY,  
student_name text,
```

```
student_city text,  
student_fees varint,  
student_phone varint  
);
```

#### **Create Data:**

```
INSERT INTO student (student_id, student_fees, student_name)  
VALUES(1,5000, 'Ajeet');
```

```
INSERT INTO student (student_id, student_fees, student_name)  
VALUES(2,3000, 'Kanchan');
```

```
INSERT INTO student (student_id, student_fees, student_name)  
VALUES(3, 2000, 'Shivani');
```

#### **READ Data**

```
SELECT * FROM student;
```

```
SELECT * FROM student WHERE student_id=2;
```

#### **Update Data:**

```
UPDATE student SET student_fees=10000,student_name='Rahul'  
WHERE student_id=2;
```

#### **DELETE Data:**

```
DELETE student_fees FROM student WHERE student_id=3;
```

## Output:

### Create, Insert and Select(Read) Query Output:

student_id	student_city	student_fees	student_name	student_phone
1	null	5000	Ajeet	null
2	null	3000	Kanchan	null
3	null	2000	Shivani	null

### Update Query Output

student_id	student_city	student_fees	student_name	student_phone
1	null	5000	Ajeet	null
2	null	10000	Rahul	null
3	null	2000	Shivani	null

(3 rows)

### Delete Query Output:

student_id	student_city	student_fees	student_name	student_phone
1	null	5000	Ajeet	null
2	null	10000	Rahul	null
3	null	null	Shivani	null

(3 rows)

## Result:

Thus the installation and configuration of Cassandra and its CURD operation are executed successfully.

Administrator: Command Prompt

Microsoft Windows [Version 10.0.19045.3930]  
(c) Microsoft Corporation. All rights reserved.

C:\Windows\system32>docker run --name cassandra-1 -d cassandra:latest

Unable to find image 'cassandra:latest' locally

latest: Pulling from library/cassandra

d66d6a6a3687: Pull complete

f0c59d8a84db: Pull complete

5e3f591e1537: Pull complete

f0374c36aa38: Pull complete

dc4b0bb2324f: Pull complete

602c5d58e79c: Pull complete

0856a9877174: Pull complete

44ee986953b1: Pull complete

06bd1714a76d: Pull complete

b6679ed26321: Pull complete

Digest: sha256:2e53494bf88d2474ad17364e3cb3b38f3a724e0a866c90b09c5e10b8f3509b61

Status: Downloaded newer image for cassandra:latest

7d9be4732ff68be11a6a5c07fec967fc4298a11c23ab433b8f1e3e1e9c3486b2

C:\Windows\system32>docker run --name cassandra-2 -d --link cassandra-1:cassandra cassandra:latest  
f245943c9e3005b0021d0ab65ccb5fff120d9fa5866041a37e0705ec18484c8c

Administrator: Command Prompt - docker exec -it cassandra-2 nodetool status

C:\Windows\system32>docker ps -a

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	
f245943c9e30	cassandra:latest	"docker-entrypoint.s..."	About a minute ago
go	Up About a minute	7000-7001/tcp, 7199/tcp, 9042/tcp, 9160/tcp	cassandra-2
7d9be4732ff6	cassandra:latest	"docker-entrypoint.s..."	About a minute ago
go	Up About a minute	7000-7001/tcp, 7199/tcp, 9042/tcp, 9160/tcp	cassandra-1
179fa94287e0	bde2020/hadoop-nodemanager:2.0.0-hadoop3.2.1-java8	"/entrypoint.sh /run..."	42 minutes ago
	Up 42 minutes (healthy)	8042/tcp	nodemanager
0ac596295dc6	bde2020/hadoop-historyserver:2.0.0-hadoop3.2.1-java8	"/entrypoint.sh /run..."	42 minutes ago
	Up 42 minutes (healthy)	8188/tcp	historyserver
eddb2b369afb	bde2020/hadoop-resourcemanager:2.0.0-hadoop3.2.1-java8	"/entrypoint.sh /run..."	42 minutes ago
	Up 42 minutes (healthy)	8088/tcp	resourcemanager
bc5f0257631c	bde2020/hadoop-namenode:2.0.0-hadoop3.2.1-java8	"/entrypoint.sh /run..."	42 minutes ago
	Up 42 minutes (healthy)	0.0.0.0:9000->9000/tcp, 0.0.0.0:9870->9870/tcp	namenode
8073557412b1	bde2020/hadoop-datanode:2.0.0-hadoop3.2.1-java8	"/entrypoint.sh /run..."	42 minutes ago
	Up 42 minutes (healthy)	9864/tcp	datanode

```

Administrator: Command Prompt
C:\Windows\system32>docker run --name cassandra-2 -d --link cassandra-1:cassandra cassandra:latest
docker: Error response from daemon: Conflict. The container name "/cassandra-2" is already in use by container "f245943c9e3005b0021d0ab65ccb5fff120d9fa5866041a37e0705ec18484c8c". You have to remove (or rename) that container to be able to reuse that name.
See 'docker run --help'.

C:\Windows\system32>docker exec -it cassandra-1 nodetool status
Datacenter: datacenter1
=====
Status=Up/Down
|/ State=Normal/Leaving/Joining/Moving
-- Address      Load      Tokens  Owns (effective)  Host ID            Rack
UN 172.17.0.4  147.92 KiB  16        100.0%          7e0aa241-5a49-4cea-ba67-eb5ff3d5fb8e  rack1
UN 172.17.0.5  158.44 KiB  16        100.0%          b1881e50-3a8f-47c8-bc4d-d2fb4b3d69d1  rack1

C:\Windows\system32>docker exec -it cassandra-2 nodetool status
Datacenter: datacenter1
=====
Status=Up/Down
|/ State=Normal/Leaving/Joining/Moving
-- Address      Load      Tokens  Owns (effective)  Host ID            Rack
UN 172.17.0.4  168.42 KiB  16        100.0%          7e0aa241-5a49-4cea-ba67-eb5ff3d5fb8e  rack1
UN 172.17.0.5  148.55 KiB  16        100.0%          b1881e50-3a8f-47c8-bc4d-d2fb4b3d69d1  rack1

```

## Node Cassandra-1:

```

Administrator: Command Prompt - docker exec -it cassandra-1 bash -c 'cqlsh'

C:\Windows\system32>docker exec -it cassandra-1 bash -c 'cqlsh'
Connected to Test Cluster at 127.0.0.1:9042
[cqlsh 6.1.0 | Cassandra 4.1.4 | CQL spec 3.4.6 | Native protocol v5]
Use HELP for help.
cqlsh> CREATE KEYSPACE test_cassandra WITH replication = {'class':'SimpleStrategy' , 'replication_factor' : 1};
cqlsh>
cqlsh> Use test_cassandra ;
cqlsh:test_cassandra> CREATE TABLE student(
    ...     student_id int PRIMARY KEY,
    ...     student_name text,
    ...     student_city text,
    ...     student_fees varint,
    ...     student_phone varint
    ... );
cqlsh:test_cassandra>
cqlsh:test_cassandra> INSERT INTO student (student_id, student_fees, student_name)
    ... VALUES(1,5000, 'Ajeet');
cqlsh:test_cassandra> INSERT INTO student (student_id, student_fees, student_name)
    ... VALUES(2,3000, 'Kanchan');
cqlsh:test_cassandra> INSERT INTO student (student_id, student_fees, student_name)
    ... VALUES(3, 2000, 'Shivani');
cqlsh:test_cassandra> SELECT * FROM student;

  student_id | student_city | student_fees | student_name | student_phone
-----+-----+-----+-----+-----+
      1 |      null |      5000 |       Ajeet |           null
      2 |      null |      3000 |      Kanchan |           null
      3 |      null |      2000 |      Shivani |           null

```

## Node Cassandra-2:

```
C:\Administrator: Command Prompt - docker exec -it cassandra-2 bash -c 'cqlsh'
Microsoft Windows [Version 10.0.19045.3930]
(c) Microsoft Corporation. All rights reserved.

C:\Windows\system32>docker exec -it cassandra-2 bash -c 'cqlsh'
Connected to Test Cluster at 127.0.0.1:9042
[cqlsh 6.1.0 | Cassandra 4.1.4 | CQL spec 3.4.6 | Native protocol v5]
Use HELP for help.
cqlsh> Use test_cassandra ;
cqlsh:test_cassandra> SELECT * FROM student;

student_id | student_city | student_fees | student_name | student_phone
-----+-----+-----+-----+-----+
  1 |      null |      5000 |     Ajeet |      null
  2 |      null |     10000 |    Rahul |      null
  3 |      null |      2000 |  Shivani |      null

(3 rows)
cqlsh:test_cassandra>
```